



CSU33031 Computer Networks

Assignment 2: Flow Forwarding

Liam Junkermann, 19300141

December 4, 2022

Contents

| | | |
|----------|------------------------------|----------|
| 1 | Introduction | 2 |
| 2 | Theory | 2 |
| 2.1 | OpenFlow | 2 |
| 3 | Implementation | 2 |
| 3.1 | Overview | 2 |
| 3.2 | Packet | 2 |
| 3.2.1 | Operation | 3 |
| 3.3 | Network Components | 3 |
| 3.3.1 | Node | 3 |
| 3.3.2 | Controller | 4 |
| 3.3.3 | Router | 4 |
| 3.3.4 | End Node | 4 |
| 4 | Future Enhancements | 5 |
| 5 | Summary | 5 |
| 6 | Reflection | 5 |

Abstract

This report will discuss the implementation of a flow forwarding software defined wide area system (SD-WAN). Starting with a brief overview of the theory needed to build and execute such a system, the components of the final implementation, a demonstration of a network topology employing this system, a brief discussion of potential improvements, and finally a summary and reflection of the assignment.

1 Introduction

This assignment focused on the development and deployment of a flow forwarding mechanism to allow two clients (End nodes) on separate networks to interact as though they have a share network by employing a SD-WAN. A centralised Controller keeps track of the Routers and in turn the End nodes attached to them. Messages are sent through an interactive console on the End node and then, with help from the Routers and Controller, directed to the appropriate End node. This report will discuss the theory which led to the final implementation.

2 Theory

This flow forwarding implementation took inspiration from the OpenFlow [1] protocol. Understanding this protocol has enabled me to lay the foundation for my implementation.

2.1 OpenFlow

This protocol was proposed as a way of allowing researchers to better test experimental protocols. Researchers found that the barrier to entry for new networking ideas was quite difficult as it is difficult to simulate the scale of a full-sized network. Therefore most new networking ideas go untested and therefore cannot make their way into production. OpenFlow addresses this by creating a standard for switches and communications which allow for programmable flow tables allowing the testing of new networking ideas. In practice OpenFlow leverages existing flow-tables built into ethernet switches and routers, this allows new protocols to be tested on a production network, while also allowing existing traffic to continue as normal.

3 Implementation

3.1 Overview

This implementation of the network uses some basic pieces from the proposed OpenFlow method. Specifically, Routers connect to the Controller with a HELLO packet, expecting a response in the form of a flow table (FWD_MOD) update packet. This table is implemented statically for the purposes of this assignment. In the virtual network for this assignment I have implemented one Controller node and five Router nodes, these are started automatically by the `docker-compose.yml` file and run headlessly for the interactive End Node nodes to connect to. There are four End Node nodes also defined in the `docker-compose.yml` file, however these processes need to be started manually.

3.2 Packet Overview

All packets have a two byte header, where byte index 0 is the packet type, and byte index one is the destination length. There are some cases where no destination is supplied, for example in the initial Router setup so this destination length packet allows data to be passed efficiently. This also allows for variable length destination names which is used in this implementation to mock different potential destinations in a humanly readable format. The packet types and their respective value are listed below:

```
// Packet Types
static final byte HELLO = 0;
static final byte PACKET_IN = 1;
static final byte FWD_MOD = 2;
static final byte NETWORK_ID = 3;
```

Listing 1: Code snippet from Node with the encoded packet types

All packets transmitted through this network follow the packet structure detailed in Figure 1




Figure 1: Standard Packet

3.2.1 Operation overview

Figure 2 Shows the steps needed for the potential message to be sent from one end node to another. The network deployed for this assignment is illustrated in Figure 3. A controller is connected to each router as well.

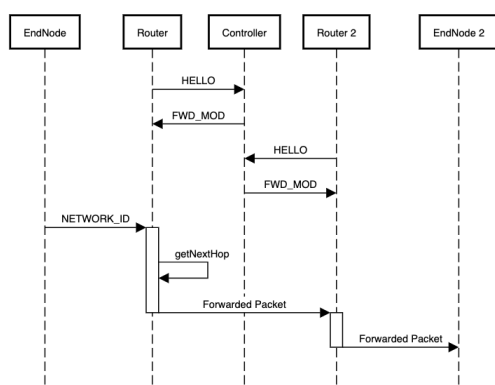


Figure 2: A sequence diagram illustrating the setup and example message sharing between two end nodes




Figure 3: An example network topology illustrating the interactions between nodes.

This implementation only requires the destination to be included in the header as the flow table manages the output of each router based on the destination, source information would therefore be redundant for the purposes of forwarding messages. This decision will be discussed further in Future Enhancements. Figure 4 and Figure 5 show the client interface and packet capture of an example interaction in the virtual network.

3.3 Network Components

3.3.1 Node

In practice each of these network nodes implement a `Node` class which handles most of the boilerplate Datagram behaviour and sets constants for header length, values, and positions, as well as providing some worker functions to generate packets and packet data appropriately. This `Node` also handles the multithreading needed to successfully multiplex within the `Controller`.

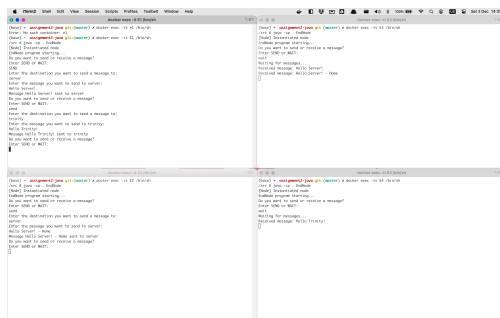


Figure 4: Four terminals showing an example of the network usage with four mocked clients

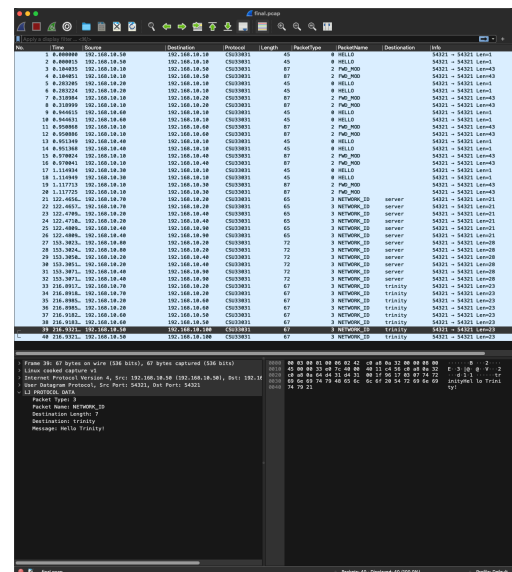


Figure 5: The packet capture of the example network usage

3.3.2 Controller

The Controller manages the forwarding table and helps routers resolve where packets should go when an entry may not be in the forwarding table. The Controller is initialised first in order to register Router nodes when receiving HELLO packets. The Controller also accepts PACKET_IN packets which are also sent from the Router. This will be discussed further in the Router section. The forwarding table has been implemented statically as two dimensional array in the current implementation of the controller. This table outlines how each Router should direct a packet based on the destination value. The Controller is a headless node and is the first node started in the Docker Compose environment.

3.3.3 Router

The Router node does most of the heavy lifting in this network, handling inputs from both End Node and other Router nodes and directing it appropriately. The Router handles two packet types:

NETWORK ID The NETWORK_ID packet is sent from the End Node nodes with messages to forward to other nodes. The Router takes the data from the incoming packet and repackages it in a new packet with the next hop destination as described in the forwarding table which is received after the Router node starts up and makes contact with the Controller

FWD MOD The FWD_MOD packet is sent from the Controller and updates the forwarding table of the Router node. This Router can also receive this packet after forwarding a packet with an unresolved destination to the controller. This will be discussed further in Future Enhancements.

3.3.4 End Node

The End Node handles the user input as well as displaying packets sent. As such, the End Node is fairly simple as it only expects to receive PACKET_IN packets with messages. In this implementation, because all input is taken from a terminal based scanner, a client must opt to WAIT for messages, a feature which will be discussed further in Future Enhancements. Ordinarily, the user will enter whether they would like to send a message or wait for messages, if they choose to send messages they will be prompted to enter the destination, and the message. This interaction is exhibited in Figure 4.

4 Future Enhancements

This project is a simplified version of what a production ready release would look like. Therefore, some decisions were made to make the work more manageable. Firstly, the static forwarding table. Because I knew what the network topology would look like for this implementation I was able to create a static forwarding table. The first enhancement was I to continue this project would be to enable Router nodes to register with one another, and the Controller to allow new nodes to be added to the WAN. End node nodes would connect to the router and be added to the forwarding table dynamically allowing for a more flexible system. Next, the user experience would need to improve allowing clients to both view messages received and sent. This is more easily done when an application can handle both user input and network traffic separately as opposed to within the same terminal. Finally I would add a source header field to allow easier packet acknowledgement and response, but this would be built based on the needs of a given application.

5 Summary

In this report discussed the requirements for the assignment, the real-world version of a solution from which I drew inspiration, and finally my implementation of a flow forwarding mechanism. Then, an exploration of the basic functionality and an example topology of the completed network, as well as some limitations and potential improvements to this implementation.

6 Reflection

Like the first assignment, this project had quite a large scope. Unfortunately during the period this assignment ran I fell somewhat behind in college work as a result of the sheer volume so I found it difficult to dedicate the time needed to execute this project with the same persistence and quality as the first assignment, this is reflected in the quality of this report. The expectations for each part of the assignment were clearer in this assignment though, which allowed me to be more organised in the final execution, and I should have used those opportunities more effectively.

References

- [1] Nick McKeown, Tom Anderson, Hari Balakrishnan, Guru Parulkar, Larry Peterson, Jennifer Rexford, Scott Shenker, and Jonathan Turner. Openflow: Enabling innovation in campus networks. *SIGCOMM Comput. Commun. Rev.*, 38(2):69–74, mar 2008.