

AMD Opteron™

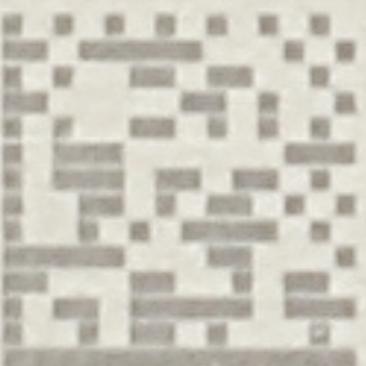
OSA146DAA5BN

CACJE 0607APBW

1418854B61344

Introduction to x64

CSU34021 - Computer Architecture II



© 2001 AMD

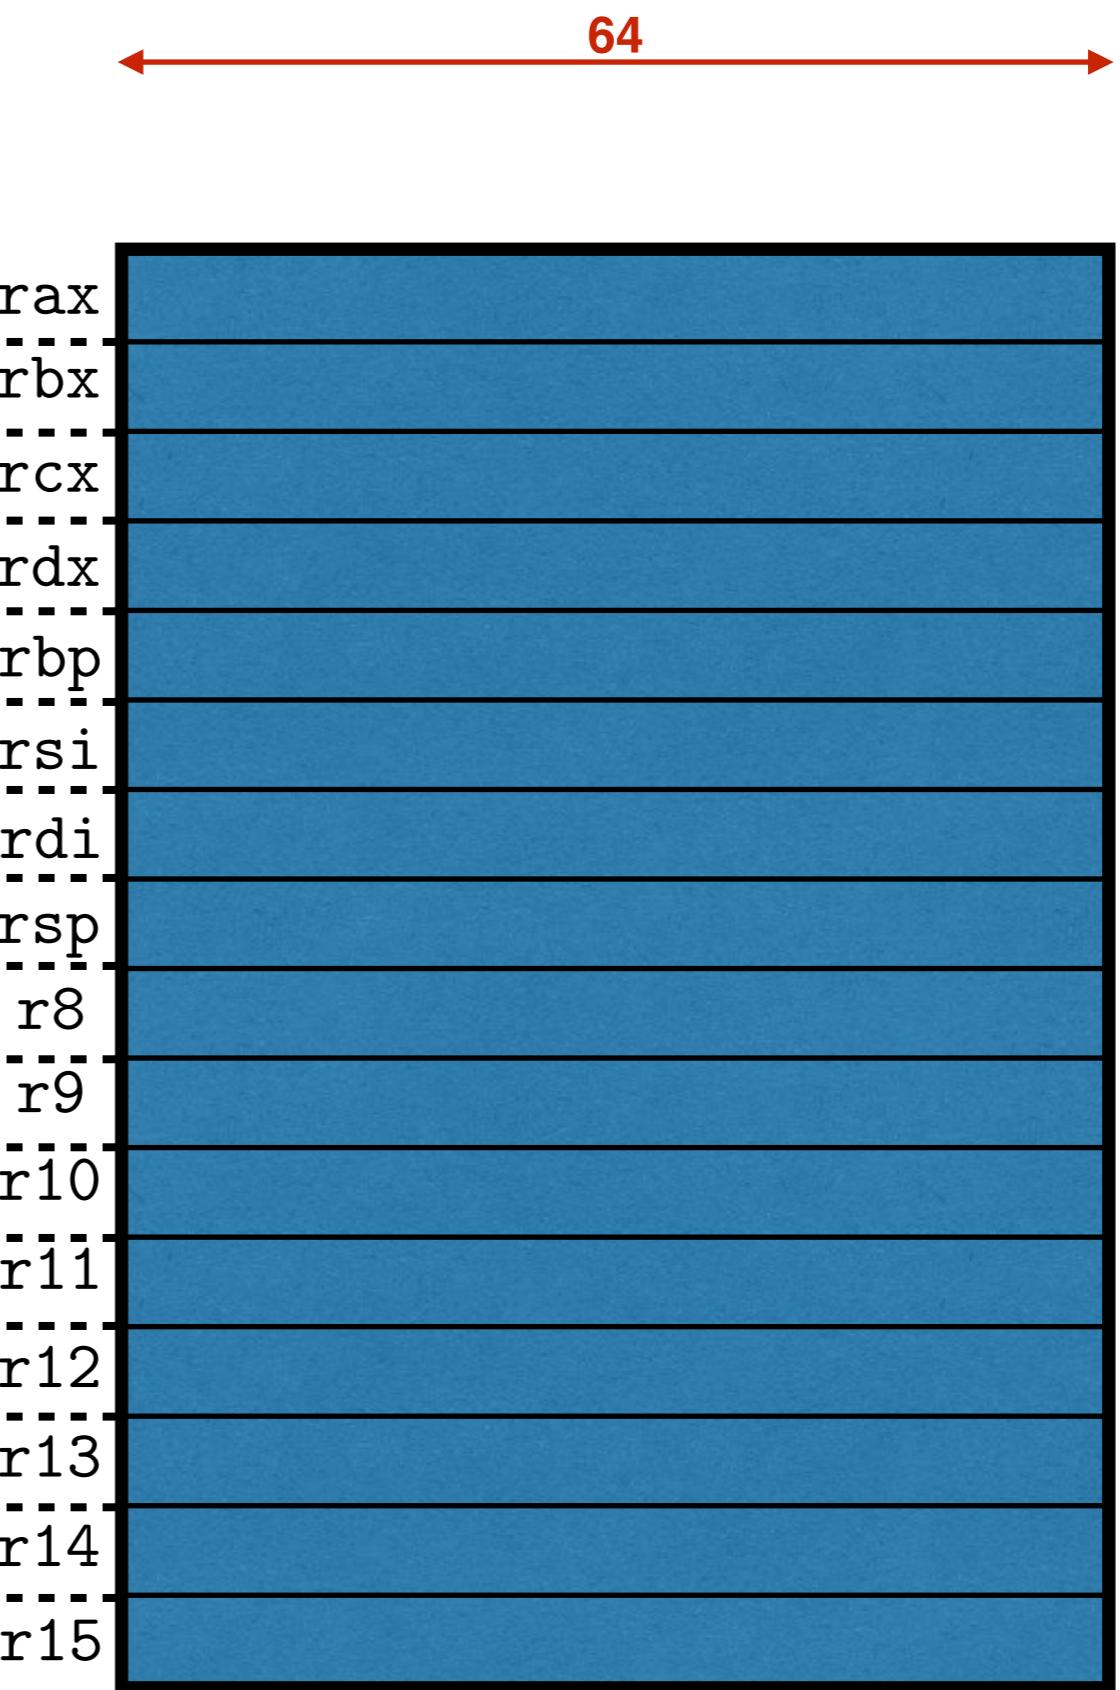
Some Background/History

- x86-64 is a backward-compatible extension of the x86 ISA.
- Originally produced by AMD (Intel had developed a completely different ISA [called IA-64] for its servers and then had to backtrack and follow up AMD footsteps).
- It has additional registers, increased memory addressing space, and the same instruction set of IA-32.



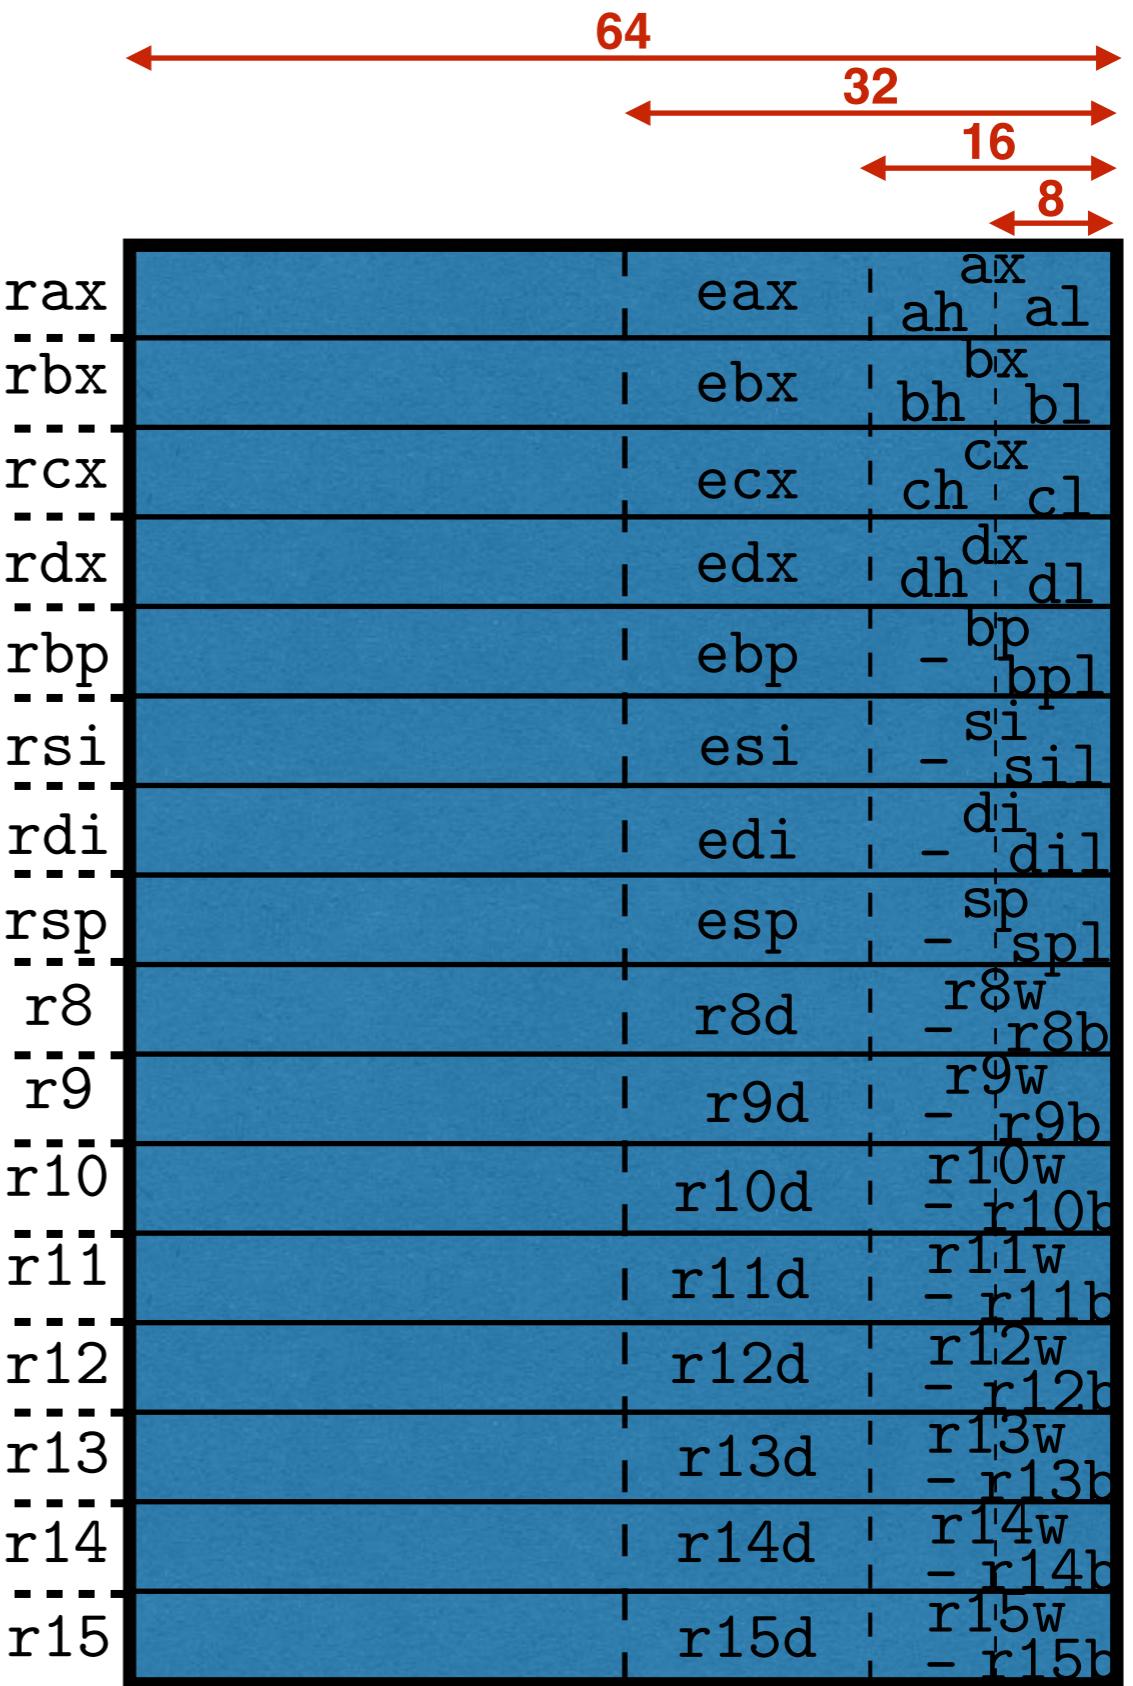
x64 Registers

- 8 more registers have been added.



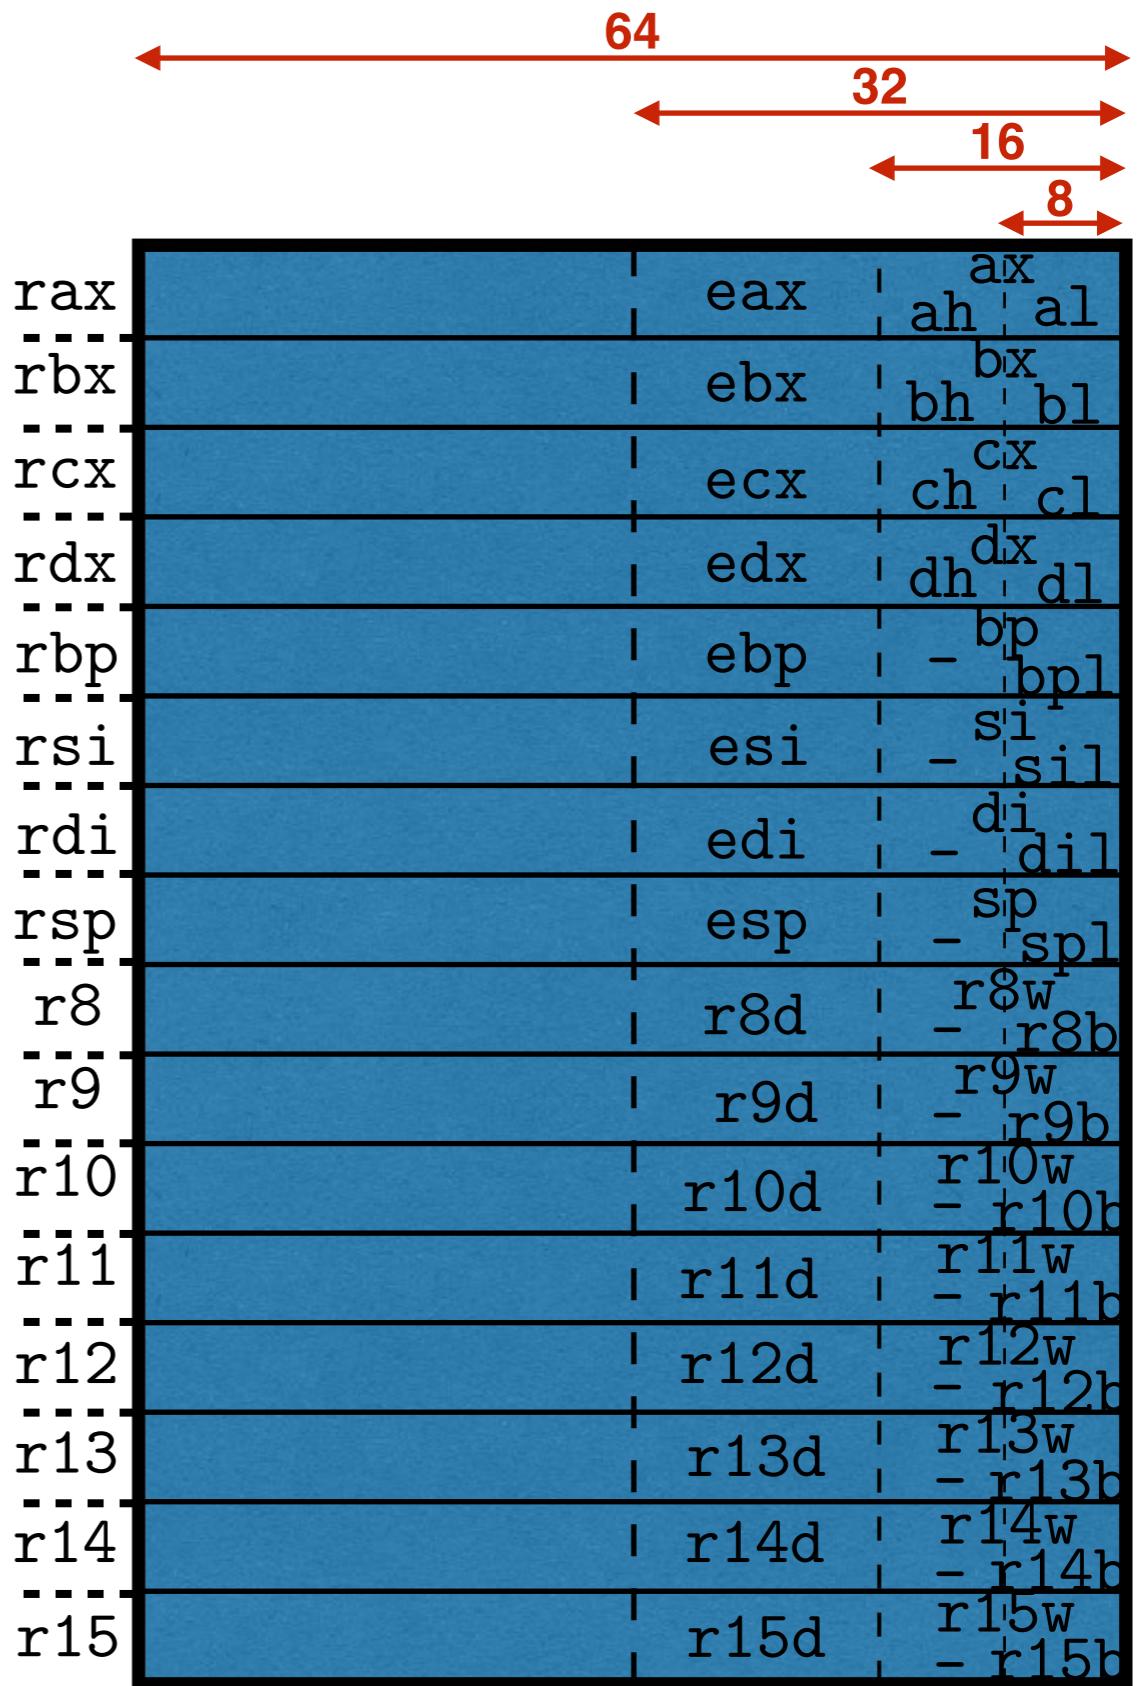
x64 Registers

- 8 more registers have been added.
- Lower bits are still accessible (not all the bits combination can be accessed at the same time...).
- All other registers extended as well.



x64 Registers

- 8 more registers have been added.
- Lower bits are still accessible (not all the bits combination can be accessed at the same time...).
- All other registers extended as well.
- Some implementations have additional registers for long integers and floating point operations (e.g. 16 XMM registers extending x87).



X64 Briefly...

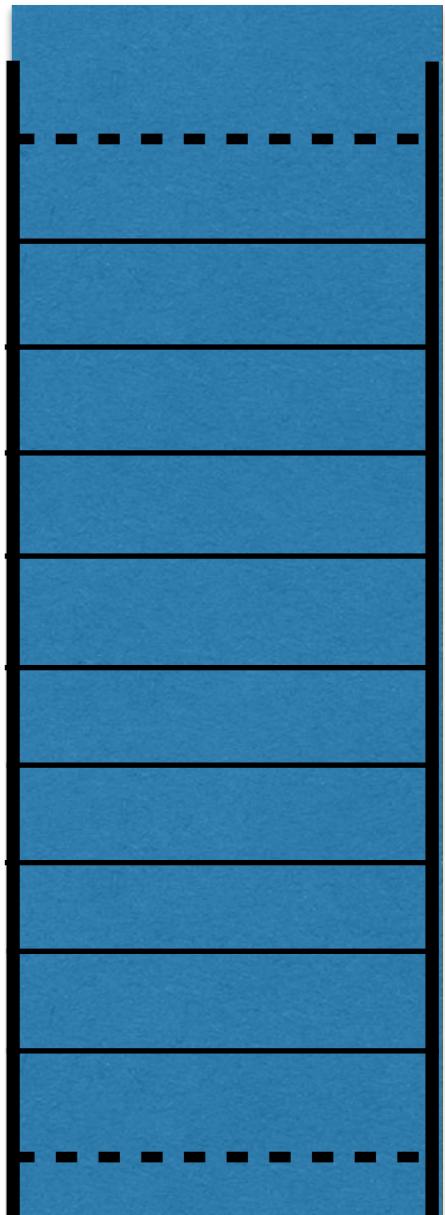
- The **instruction set** is virtually the **same** of x86.
- Main difference in the **calling convention** for functions. Let's see the Microsoft one...
- The calling convention uses the fact that there are **more registers** to **speed up** the function calling.
- Plus, having more registers at ones disposal, often the **frame pointer** for the stack is not even used.

Microsoft Calling Convention

- `rcx`, `rdx`, `r8`, `r9` used for the first 4 integer or pointer arguments (**from left to right**).
- Additional parameters pushed on the stack (**right to left**).
- Caller to allocate 32-bytes of **shadow space** for storing the input registers. Caller also has to clean the stack.
- Integers returned to `rax`. For larger return values pointers are used.
- `rax`, `rcx`, `rdx`, `r8`, `r9`, `r10`, `r11` **considered volatile**.
- If other registers are used, the callee needs to preserve them.
- **The stack needs to be aligned on a 16-byte boundary!**

```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)  
{_int64 x; _int64 y; ... return x;}
```

rsp →



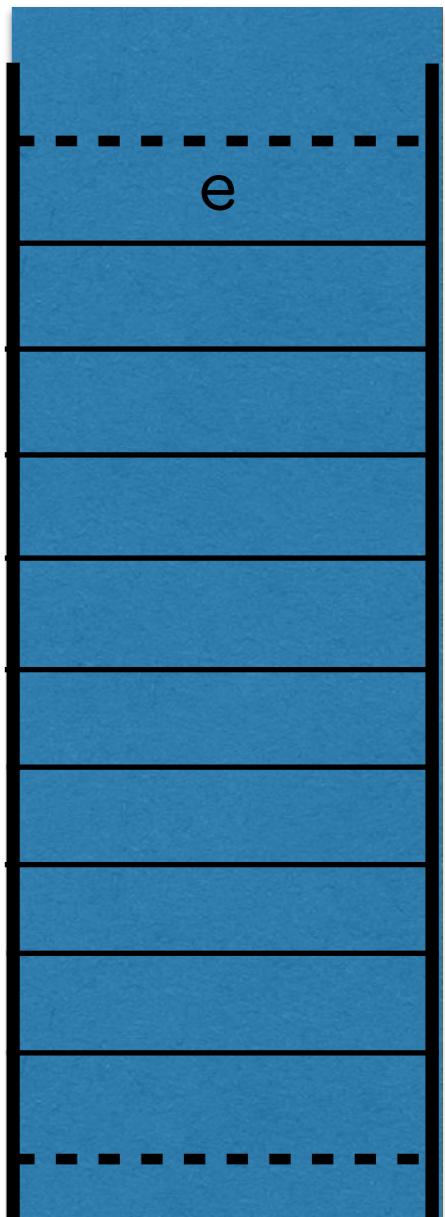
```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)
```

```
{_int64 x; _int64 y; ... return x;}
```

push e

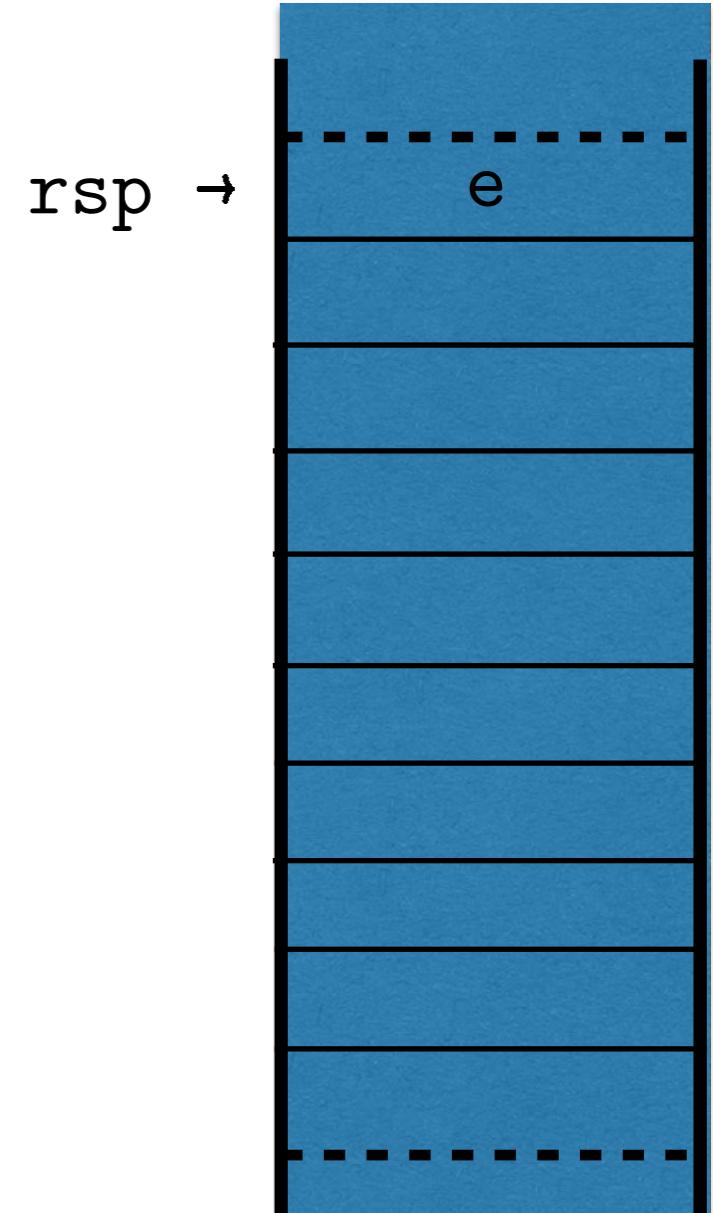
;5th arg in stack

rsp →



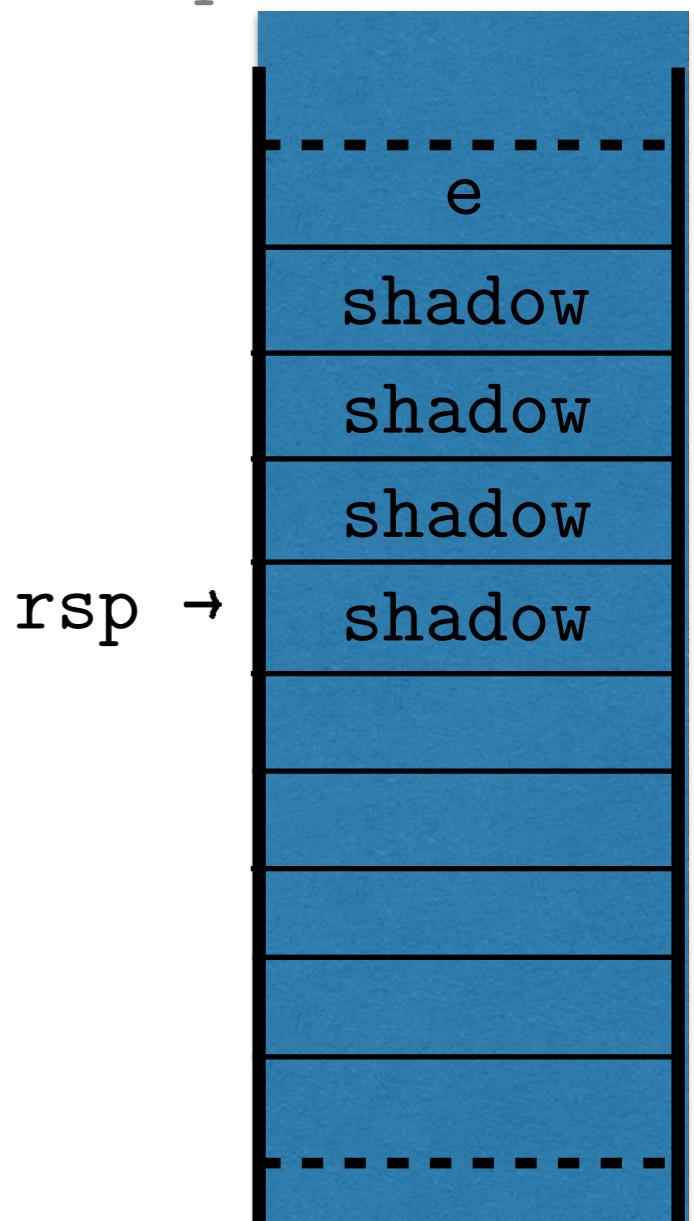
```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)  
{_int64 x; _int64 y; ... return x;}
```

```
push e ;5th arg in stack  
mov rcx, a ;1st arg  
mov rdx, b ;2nd arg  
mov r8 , c ;3rd arg  
mov r9 , d ;4th arg
```



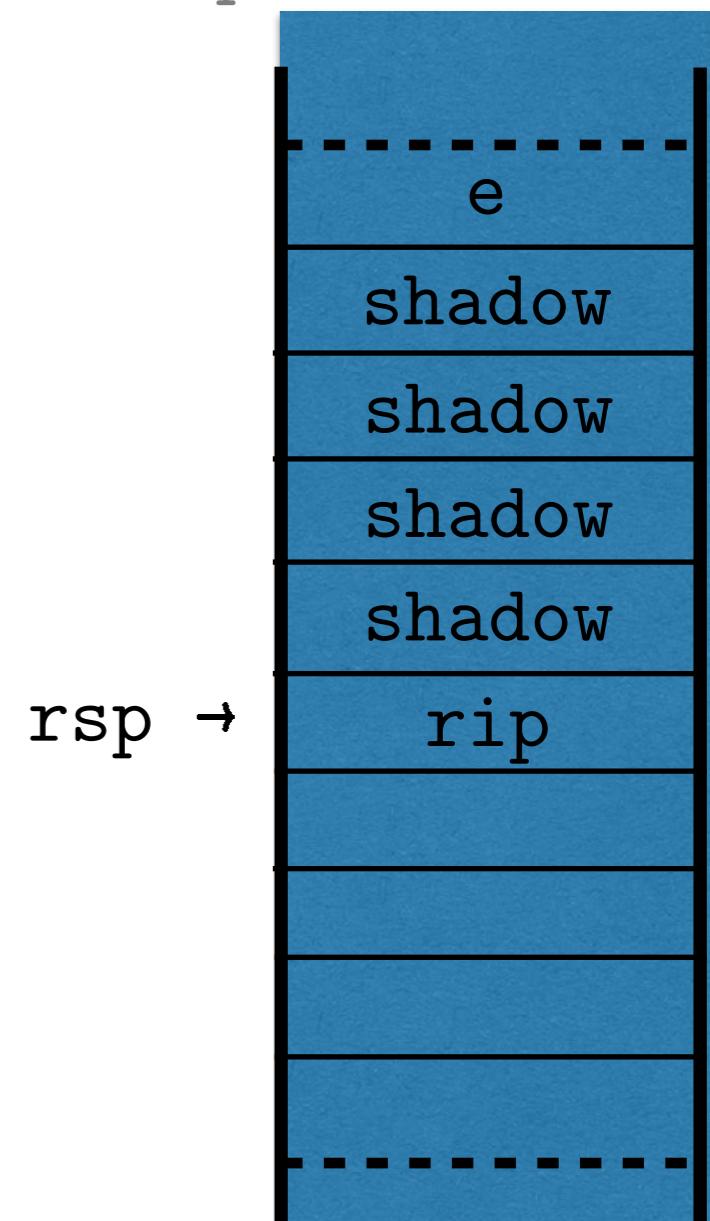
```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)  
{_int64 x; _int64 y; ... return x;}
```

```
push e ;5th arg in stack  
mov rcx, a ;1st arg  
mov rdx, b ;2nd arg  
mov r8 , c ;3rd arg  
mov r9 , d ;4th arg  
sub rsp, 32 ;shadow space
```



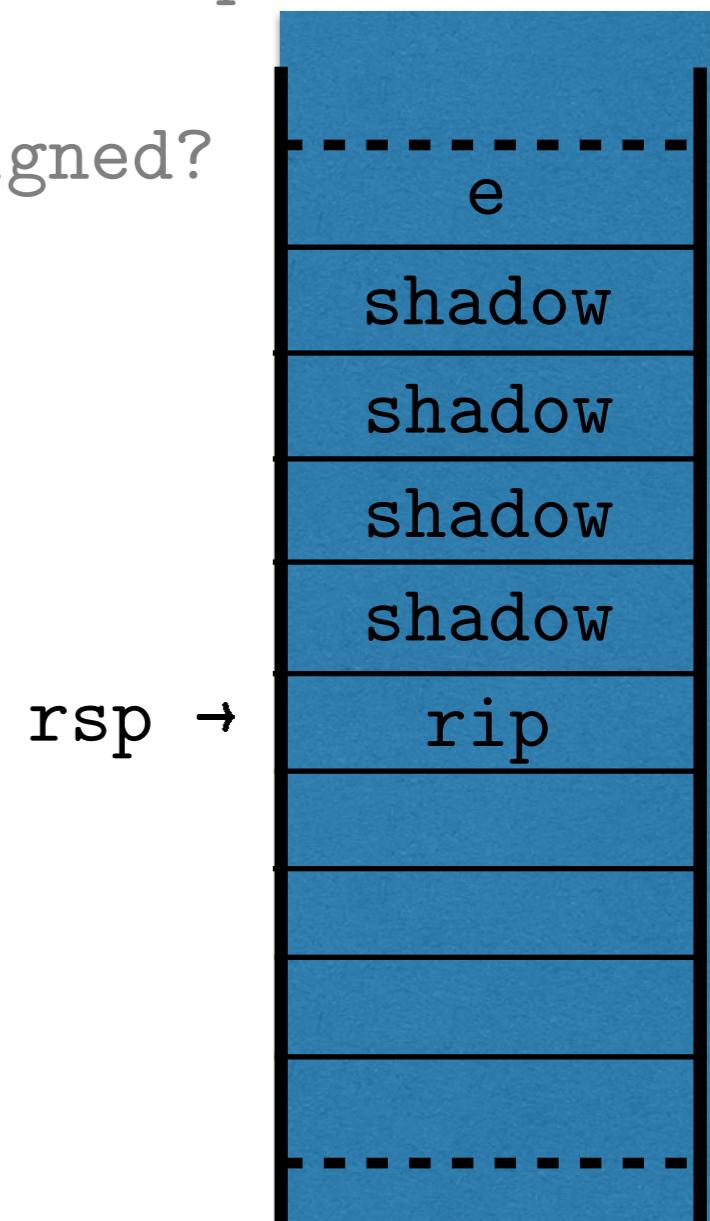
```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)  
{_int64 x; _int64 y; ... return x;}
```

```
push e ;5th arg in stack  
mov rcx, a ;1st arg  
mov rdx, b ;2nd arg  
mov r8 , c ;3rd arg  
mov r9 , d ;4th arg  
sub rsp, 32 ;shadow space  
call my_fun
```



```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)  
{_int64 x; _int64 y; ... return x;}
```

```
push e ;5th arg in stack  
mov rcx, a ;1st arg  
mov rdx, b ;2nd arg  
mov r8 , c ;3rd arg  
mov r9 , d ;4th arg  
sub rsp, 32 ;shadow space  
call my_fun  
;is the stack aligned?
```



```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)
{_int64 x; _int64 y; ... return x;}
```

my_fun PROC:

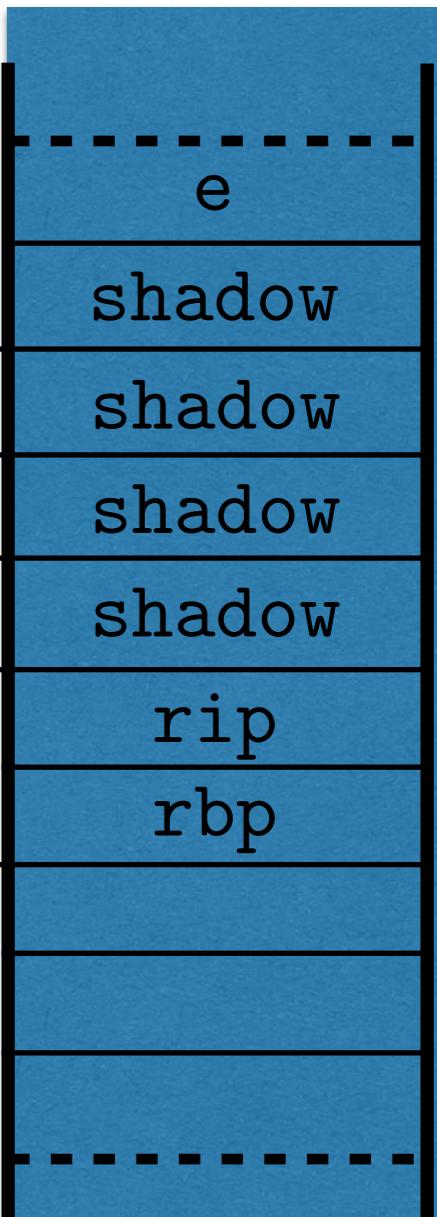
;creating stack frame

```
push rbp
```

```
mov rbp, rsp
```

```
push e          ;5th arg in stack
mov rcx, a     ;1st arg
mov rdx, b     ;2nd arg
mov r8 , c     ;3rd arg
mov r9 , d     ;4th arg
sub rsp, 32    ;shadow space
call my_fun
;is the stack aligned?
```

rsp → rbp →

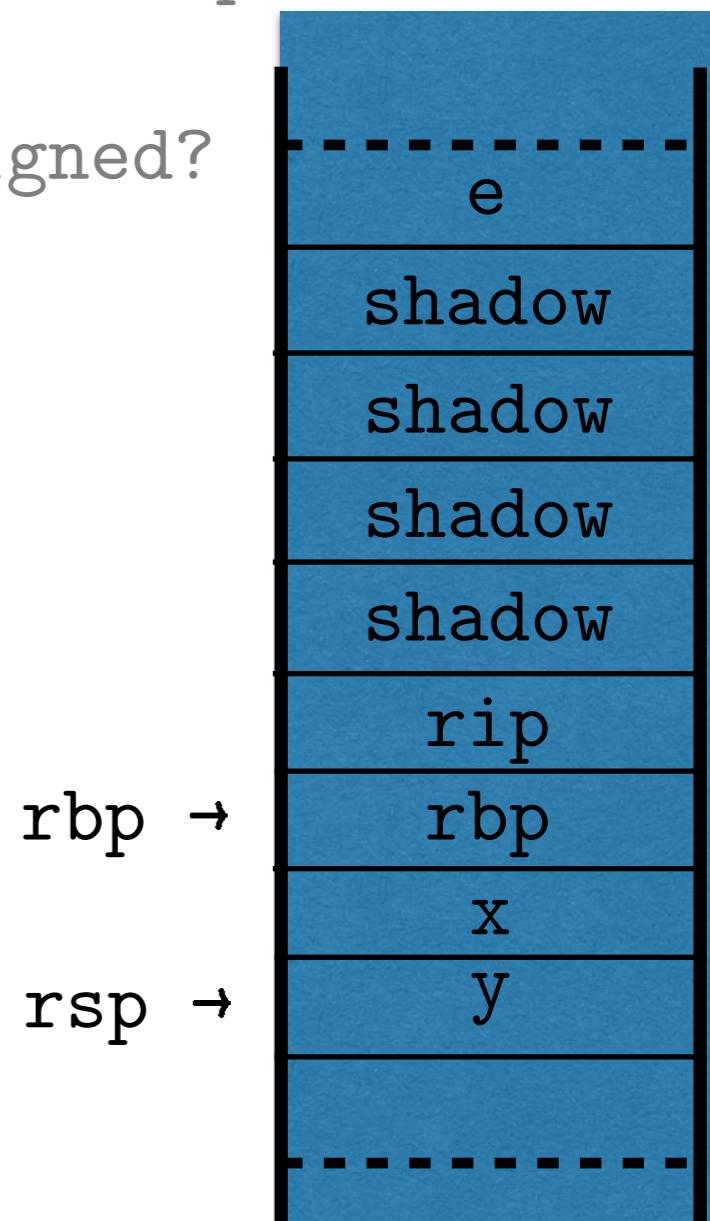


```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)
{_int64 x; _int64 y; ... return x;}
```

my_fun PROC:

```
;creating stack frame
push rbp
mov rbp, rsp
;local variables
sub rsp, 16
```

```
push e          ;5th arg in stack
mov rcx, a     ;1st arg
mov rdx, b     ;2nd arg
mov r8 , c     ;3rd arg
mov r9 , d     ;4th arg
sub rsp, 32    ;shadow space
call my_fun
;is the stack aligned?
```

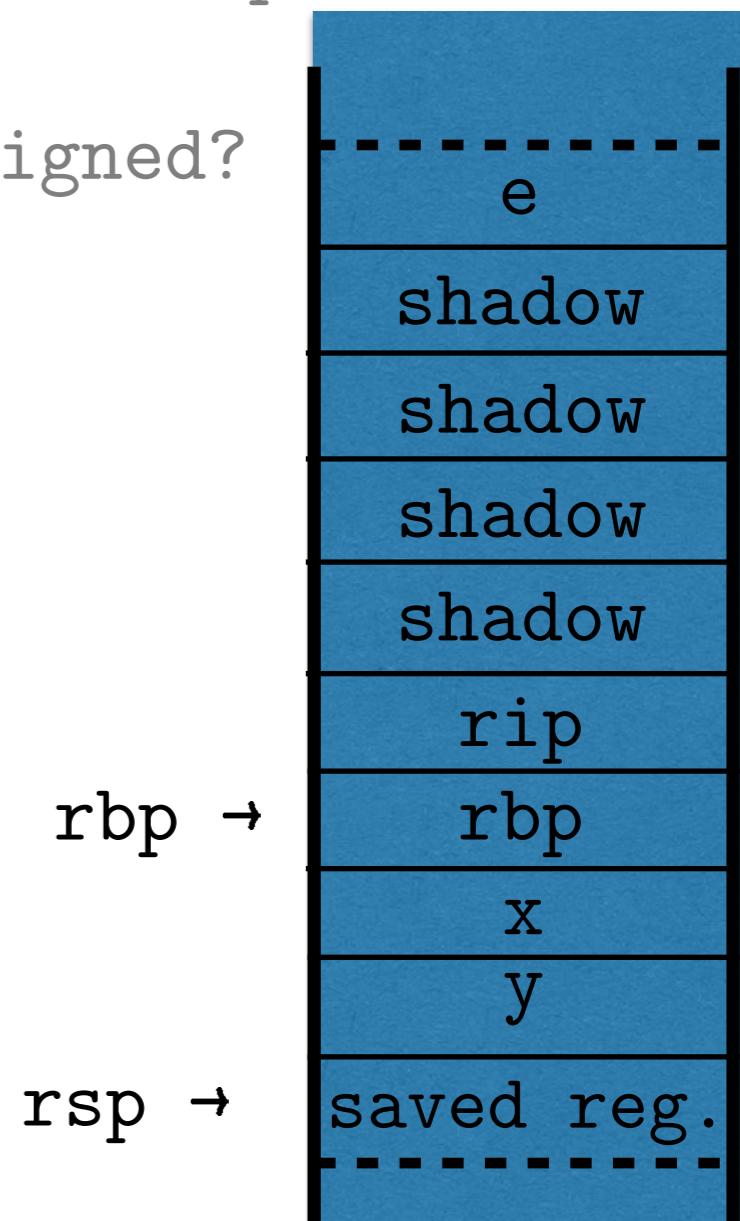


```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)
{_int64 x; _int64 y; ... return x;}
```

my_fun PROC:

```
;creating stack frame
push rbp
mov rbp, rsp
;local variables
sub rsp, 16
;saving preserved regs
;(if necessary)
push rbx
...
...
```

```
push e ;5th arg in stack
mov rcx, a ;1st arg
mov rdx, b ;2nd arg
mov r8 , c ;3rd arg
mov r9 , d ;4th arg
sub rsp, 32 ;shadow space
call my_fun
;is the stack aligned?
add rsp, 40
```



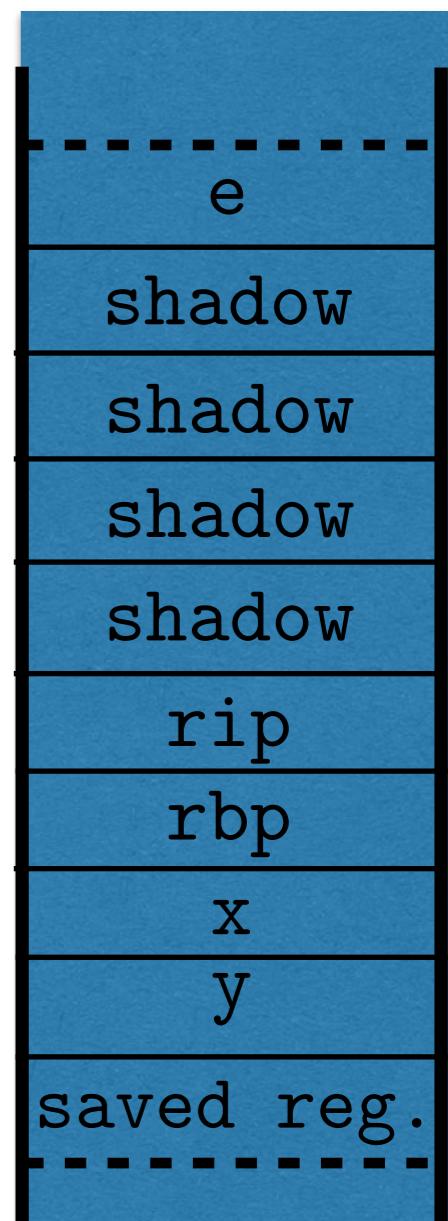
```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)  
{_int64 x; _int64 y; ... return x;}
```

my_fun PROC:

```
;creating stack frame  
push rbp  
mov rbp, rsp  
;local variables  
sub rsp, 16  
;saving preserved regs  
;(if necessary)  
push rbx  
...  
;cleaning the stack  
pop rbx  
add rsp, 16  
pop rbp
```

```
push e ;5th arg in stack  
mov rcx, a ;1st arg  
mov rdx, b ;2nd arg  
mov r8 , c ;3rd arg  
mov r9 , d ;4th arg  
sub rsp, 32 ;shadow space  
call my_fun  
;is the stack aligned?  
add rsp, 40
```

rsp →



```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)
```

```
{_int64 x; _int64 y; ... return x;}
```

```
my_fun PROC:
```

```
;creating stack frame
```

```
push rbp
```

```
mov rbp, rsp
```

```
;local variables
```

```
sub rsp, 16
```

```
;saving preserved regs
```

```
;(if necessary)
```

```
push rbx
```

```
...
```

```
;cleaning the stack
```

```
pop rbx
```

```
add rsp, 16
```

```
pop rbp
```

```
;returning
```

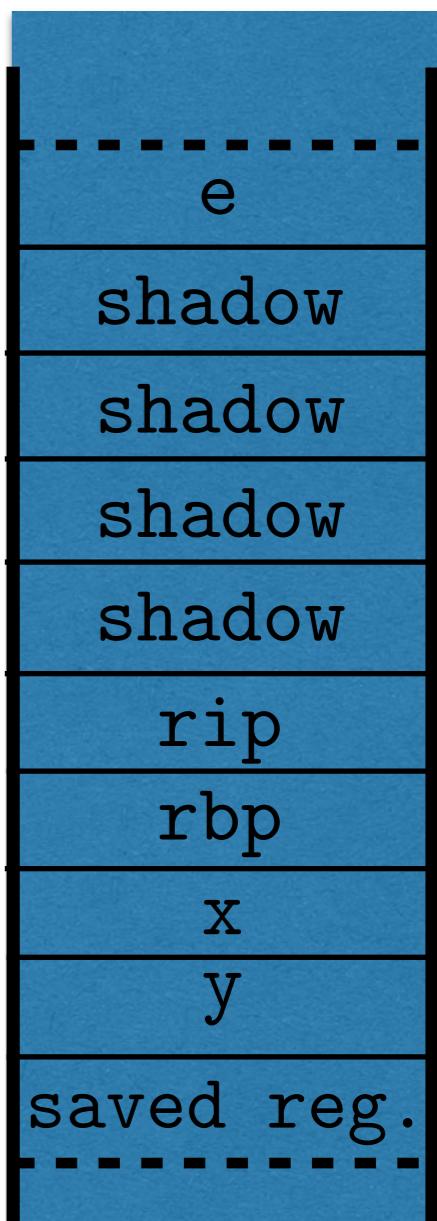
```
mov rax, [rbp+8]
```

```
ret 0
```

```
my_fun ENDP
```

```
push e ;5th arg in stack
mov rcx, a ;1st arg
mov rdx, b ;2nd arg
mov r8 , c ;3rd arg
mov r9 , d ;4th arg
sub rsp, 32 ;shadow space
call my_fun
;is the stack aligned?
```

rsp →



```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)
```

```
{_int64 x; _int64 y; ... return x;}
```

```
my_fun PROC:
```

```
;creating stack frame
```

```
push rbp
```

```
mov rbp, rsp
```

```
;local variables
```

```
sub rsp, 16
```

```
;saving preserved regs
```

```
;(if necessary)
```

```
push rbx
```

```
...
```

```
;cleaning the stack
```

```
pop rbx
```

```
add rsp, 16
```

```
pop rbp
```

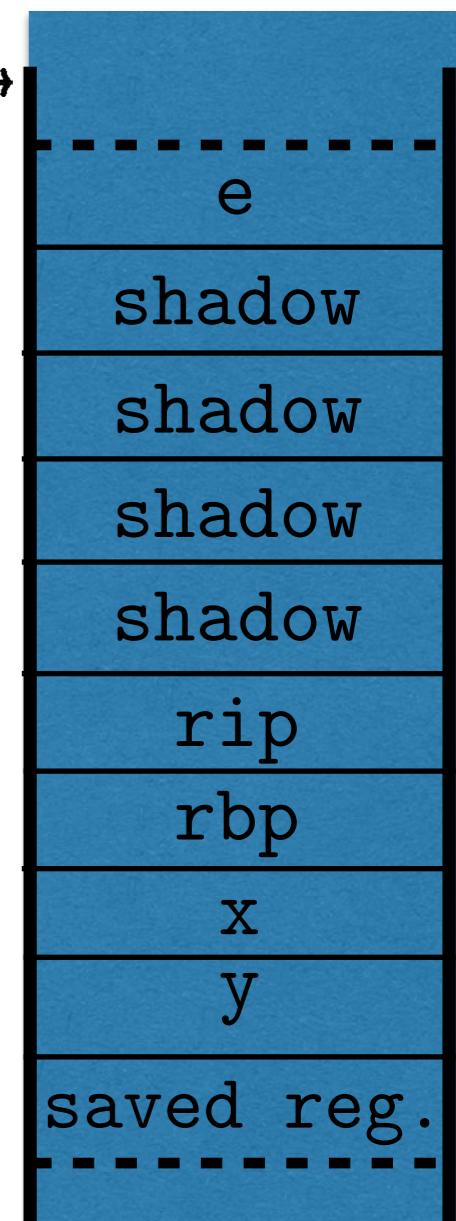
```
;returning
```

```
mov rax, [rbp+8]
```

```
ret 0
```

```
my_fun ENDP
```

```
push e ;5th arg in stack
mov rcx, a ;1st arg
mov rdx, b ;2nd arg
mov r8 , c ;3rd arg
mov r9 , d ;4th arg
sub rsp, 32 ;shadow space
call my_fun
add rsp, 40
```



```
_int64 my_fun(_int64 a, _int64 b, _int64 c, _int64 d, _int64 e)
```

```
{_int64 x; _int64 y; ... return x;}
```

```
my_fun PROC:
```

```
;creating stack frame
```

```
push rbp
```

```
mov rbp, rsp
```

```
;local variables
```

```
sub rsp, 16
```

```
;saving preserved regs
```

```
;(if necessary)
```

```
push rbx
```

```
...
```

```
;cleaning the stack
```

```
pop rbx
```

```
add rsp, 16
```

```
pop rbp
```

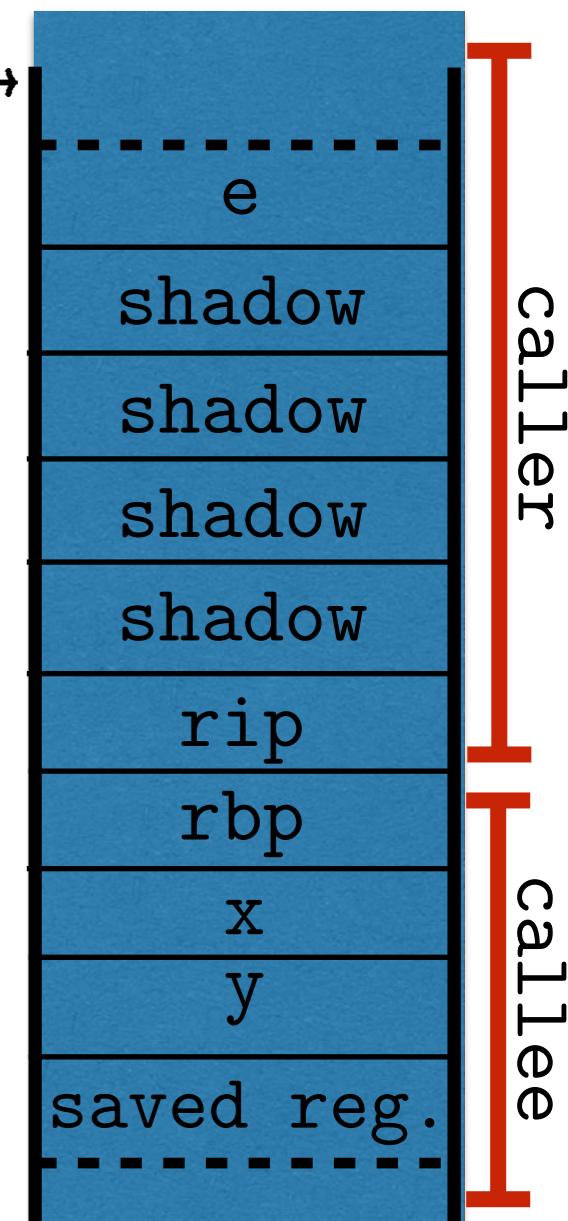
```
;returning
```

```
mov rax, [rbp+8]
```

```
ret 0
```

```
my_fun ENDP
```

```
push e ;5th arg in stack
mov rcx, a ;1st arg
mov rdx, b ;2nd arg
mov r8 , c ;3rd arg
mov r9 , d ;4th arg
sub rsp, 32 ;shadow space
call my_fun
add rsp, 40
```



Calling Convention and Shadow Space

- Advantage of `x64` Microsoft Calling Convention vs. the `x86` one?
- The **full shadow space** needs to be allocated **independently of number of input parameters**.
- It is used for `debugging`, for `spilling input registers into the stack` (useful for varargs functions), and for `scratch space` for the **callee function**.
- Shadow space doesn't need to be recreated for sequential calls!

Shadow Space

```
sub    rsp      , 48          int my_fun()  
                                ...  
                                int main(){  
                                ...  
                                my_fun(a,b,c,d,e)  
                                my_fun(a,b,c,d,f,g)  
                                ...  
                                }
```

Shadow Space

```
sub    rsp      , 48 52          int my_fun()  
                                ...  
int main(){  
    ...  
    my_fun(a,b,c,d,e)  
    my_fun(a,b,c,d,f,g)  
    ...  
}
```

Shadow Space

```
sub    rsp      , 48 52          int my_fun()  
mov    rax      , e  
mov    [rsp+40], rax ;5th arg in stack  
mov    rcx      , a ;1st arg  
mov    rdx      , b ;2nd arg  
mov    r8       , c ;3rd arg  
mov    r9       , d ;4th arg  
call   my_fun  
  
int main(){  
...  
my_fun(a,b,c,d,e)  
my_fun(a,b,c,d,f,g)  
...  
}
```

Shadow Space

```
sub    rsp      , 48 52          int my_fun()  
mov    rax      , e  
mov    [rsp+40], rax ;5th arg in stack  
mov    rcx      , a ;1st arg  
mov    rdx      , b ;2nd arg  
mov    r8       , c ;3rd arg  
mov    r9       , d ;4th arg  
call   my_fun  
  
mov    rax      , g  
mov    [rsp+48], rax ;6th arg in stack  
mov    rax      , f  
mov    [rsp+40], rax ;5th arg in stack  
call   my_fun
```

```
...  
int main(){  
...  
my_fun(a,b,c,d,e)  
my_fun(a,b,c,d,f,g)  
...  
}
```

Shadow Space

```
sub    rsp      , 48 52          int my_fun()  
mov    rax      , e  
mov    [rsp+40], rax ;5th arg in stack  
mov    rcx      , a ;1st arg  
mov    rdx      , b ;2nd arg  
mov    r8       , c ;3rd arg  
mov    r9       , d ;4th arg  
call   my_fun  
  
mov    rax      , g  
mov    [rsp+48], rax ;6th arg in stack  
mov    rax      , f  
mov    [rsp+40], rax ;5th arg in stack  
call   my_fun  
add   rsp      , 52
```

```
...  
int main(){  
...  
my_fun(a,b,c,d,e)  
my_fun(a,b,c,d,f,g)  
...  
}
```

Code Generation Strategy for Non-leaf function x64

- Allocate enough stack space to support the call to the function with the most parameters [ONCE at the start of the function].
- Use the same stack space [and register, ofc] to pass parameters to ALL the functions called [Don't push, but mov!]
- Be careful that the stack is 16-byte aligned upon entering a function, otherwise it could crash!

Summary

- Details on x86 Intel architecture.
- Introduction to x86 and x64 ISA and MASM assembler.
- Main conventions for function calling in Microsoft.
- Where next?
- x86 and x64 probably include the majority of personal devices used in real life. There is a plethora of information available on the web, and vibrant communities around them.
- If you want to read more:
 - “Introduction to Assembly language programming” (Dandamudi).
 - “The Intel 64 and IA-32 Architecture Software Developer Manuals”.